

Introduction to Data Mining: Assignment #3

Summer 2017

Due: June 15th, 23:59:59 CST (UTC +8).

1. Neural Networks

The codes of this section are in the *neural_networks* folder.

In this problem, we will implement the feedforward and backpropagation process of the neural networks. We will use *digital.mat* as our experiment data. Finish *fullyconnect_feedforward.m*, *fullyconnect_backprop.m*, *relu_feedforward.m*, *relu_backprop.m* and the testing part in *run.m*. Then we can train three layer (data, hidden-relu, loss) neural networks and report test accuracy.

Supplementary Knowledges:

- i) Instead of using MSE loss function, we adopt softmax loss function here. Recall that in Assignment #2, we used logistic regression to classify two classes. And softmax regression model is a model that extends logistic regression to classify more classes than two¹. In this problem, we have 10 classes. The softmax loss part codes are done.
- ii) We can use *gradient_check.m* to check the correctness your computation. If $\frac{d}{d\theta}J(\theta) = \lim_{\epsilon \rightarrow 0} \frac{J(\theta+\epsilon) - J(\theta)}{\epsilon}$ holds, then we are in the right way.²
- iii) Weight decay and momentum are used to update weight parameters in *get_new_weight_inc.m*.

2. K-Nearest Neighbor

The codes of this section are in the *knn* folder.

In this problem, we will play with K-Nearest Neighbor (KNN) algorithm and try it on real-world data. Implement KNN algorithm (in *knn.m*), then answer the following questions.

- (a) In *knn_exp.m*, try KNN with different K (you should at least experiment K = 1, 10 and 100) and plot the decision boundary.

You are encouraged to vectorize³ your code, otherwise the experiment time might be extremely long. You may find the MATLAB build-in functions *pdist2*, *sort*, *max* and *hist* useful. Also, you can use the function *eudist2*⁴ written by Prof. Deng Cai⁵.

- (b) We have seen the effects of different choices of K. How can you choose a proper K when dealing with real-world data ?
- (c) Now let us use KNN algorithm to hack the CAPTCHA of a website⁶ that we are all

¹http://ufldl.stanford.edu/wiki/index.php/Softmax_Regression

²http://ufldl.stanford.edu/wiki/index.php/Gradient_checking_and_advanced_optimization

³http://www.mathworks.cn/cn/help/matlab/matlab_prog/vectorization.html

⁴<http://www.cad.zju.edu.cn/home/dengcai/Data/code/EuDist2.m>

⁵Prof. Deng Cai is an expert on MATLAB, you can find all his code at <http://www.cad.zju.edu.cn/home/dengcai/Data/data.html>. You can learn how to write fast MATLAB code by reading his code.

⁶<http://jwbinfosys.zju.edu.cn/default2.aspx>

familiar with:



Finish *hack.m* to recognize the CAPTCHA image using KNN algorithm.

You should label some training data yourself, and store the training data in *hack_data.mat*. Helper functions *extract_image* and *show_image* are give for your convenience.

Remember to submit *hack_data.mat* along with your code and report.

3. Decision Tree and ID3

Consider the scholarship evaluation problem: selecting scholarship recipients based on gender and GPA. Given the following training data:

Gender	GPA	Scholarship	Count
F	Low	+	10
F	High	+	95
M	Low	+	5
M	High	+	90
F	Low	-	80
F	High	-	20
M	Low	-	120
M	High	-	30

Draw the decision tree that would be learned by ID3 algorithm and annotate each non-leaf node in the tree with the information gain attained by the respective split.

4. K-Means Clustering

The codes of this section are in the *knn* folder.

Finally, we will run our first unsupervised algorithm – k-means clustering. Implement k-means algorithm (in *kmeans.m*), then answer the following questions.

Note that there are different kind of methods to setup initial cluster centers for k-means algorithm, we will use a simple one – randomly choose K samples from dataset as initial cluster centers.

- (a) Run your k-means algorithm on *kmeans_data.mat* with the number of clusters K set to 2. Repeat the experiment 1000 times. Use *kmeans_plot.m* to visualize the process of k-means algorithm for the two trials with largest and smallest SD (sum of distances from each point to its respective centroid).
- (b) You should observe the issue that the outcome of k-means algorithm is very sensitive to cluster centroids initialization from the above experiment. How can we get a stable result using k-means?
- (c) Run your k-means algorithm on the digit dataset *digit_data.mat* with the number of clusters K set to 10, 20 and 50. Visualize the centroids using *show_digit.m*. You should be able to observe that k-means algorithm can discover the patterns in dataset without any label information.
- (d) Another important application of k-means is Vector quantization⁷. Vector quantization is a classical quantization technique from signal processing. It works by dividing a large set of points (vectors) into groups, then representing the data points by their group centroid points, as in k-means and some other clustering algorithms.

Here we will use vector quantization to do image compression. By clustering image pixel value into K groups, we can represent each pixel with $\log(K)$ bits, instead of 24 bits (RGB, each channel has 8bit depth).

Finish *vq.m*. Compress images with K set to 8, 16, 32 and 64. I have provided you some sample images, however use your own photos is encouraged.

What is the compress ratio if we set K to 64 (Optionally, you can compute the compress ratio using Huffman encoding) ?

Please submit your homework report to at <http://assignment.zjulearning.org:8081/> in pdf format, with all your code in a zip archive.

⁷https://en.wikipedia.org/wiki/Vector_quantization